

Probabilistic Graph Models: from Bayesian to Factor Graphs

ECE/CS 498 DS U/G

Lecture 13

Ravi K. Iyer

Dept. of Electrical and Computer Engineering
University of Illinois at Urbana Champaign

Announcements

- Today:
 - Monday March 4
 - Hidden Markov Model
 - Factor Graphs
- Wednesday (March 6):
 - Solve last year's midterm
- HW3 has been released; due on Friday, March 8
- MP2 Checkpoint 2 due on Friday, March 8
- Midterm on Monday (March 11) (**Note: Room change**)
 - ECE room 2013 (NetID starting with **a-m**), room 3013 (NetID starting with **n-z**)
 - Will start on time at 12:30 – 1:50pm
 - Bring one 8x11 sheet of notes
 - Closed book, no calculators or electronic devices
 - Bring your Univ IDs

Announcements

CSL Special Seminar

Join us Monday, March 4th, 4:00 p.m., CSL B02
(Followed by a Pizza Social hosted by Anthem)

Anthem, Inc. **The Practical Application of Data Science in Health Insurance**



Shawn Wang
VP, Data Science at Anthem

- 10+ years of experience leading advanced analytic teams for both retail and health insurance industries.
- Currently leading the Anthem Data Science team in developing and enabling industry leading data science artificial-intelligence solutions
- Master's degree in Information Management from the University of Southern California's Marshall School of Management

Email: shawn.wang@anthem.com

LinkedIn: <https://www.linkedin.com/in/shawn-wang-47164612/>



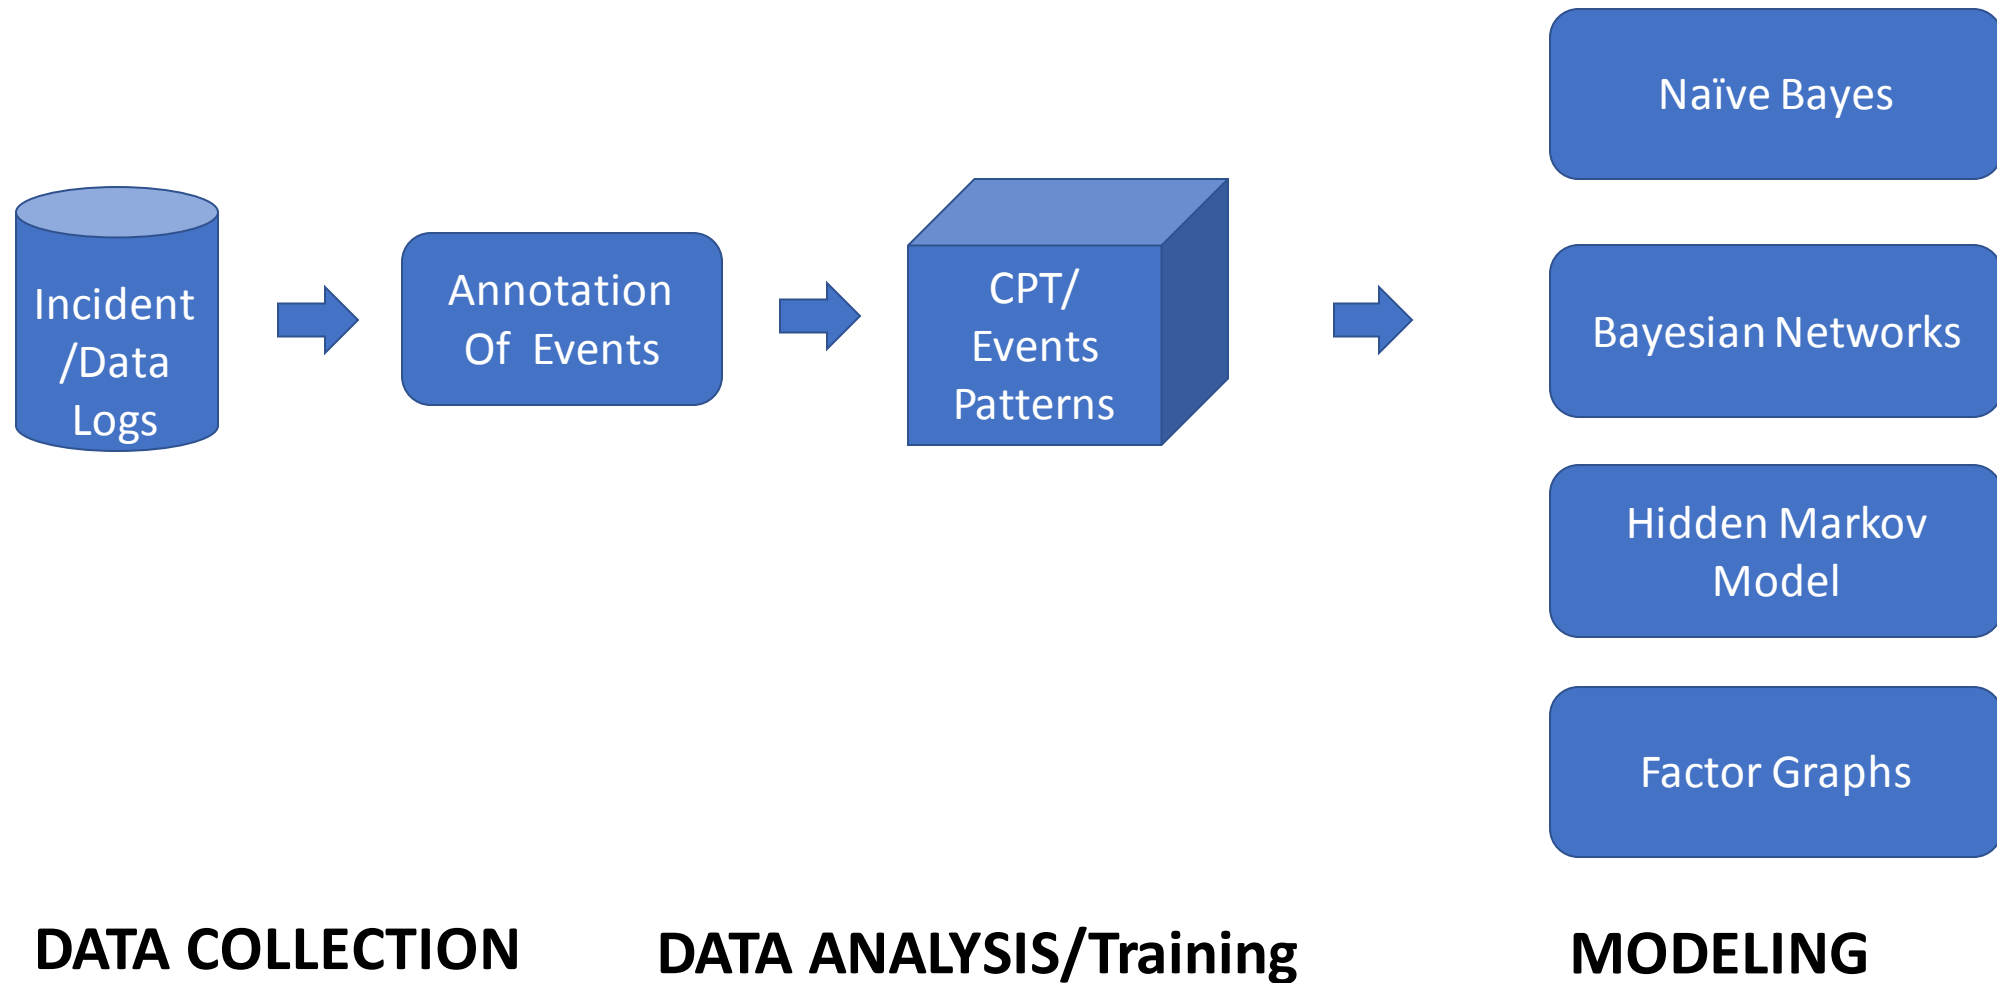
Adarsh Ramesh
Staff VP, Advanced Analytics

- 10+ years of experience leading advanced analytic teams in CPG, Retail and health insurance industries.
- Currently leading the Data Science Machine Learning team supporting solution development across various Anthem functions, including risk adjustment, program integrity, provider and consumer analytics, and more
- MBA degree from Penn State University and Undergraduate degree in Electronics Engineering

Email: adarsh.ramesh@anthem.com

LinkedIn: <https://www.linkedin.com/in/aramesh/>

Overview of PGM Data Analytics/Modeling Process





Measurements from NCSA@Illinois: Five minute Snap Shot

- Goals:


- Provide a system-level characterization of incidents and evaluate the intricacies of real-time diagnosis
- Design protection strategies to reduce missed incidents and false positives
- Experimentally Demonstrate new techniques in a sandbox

- Challenges

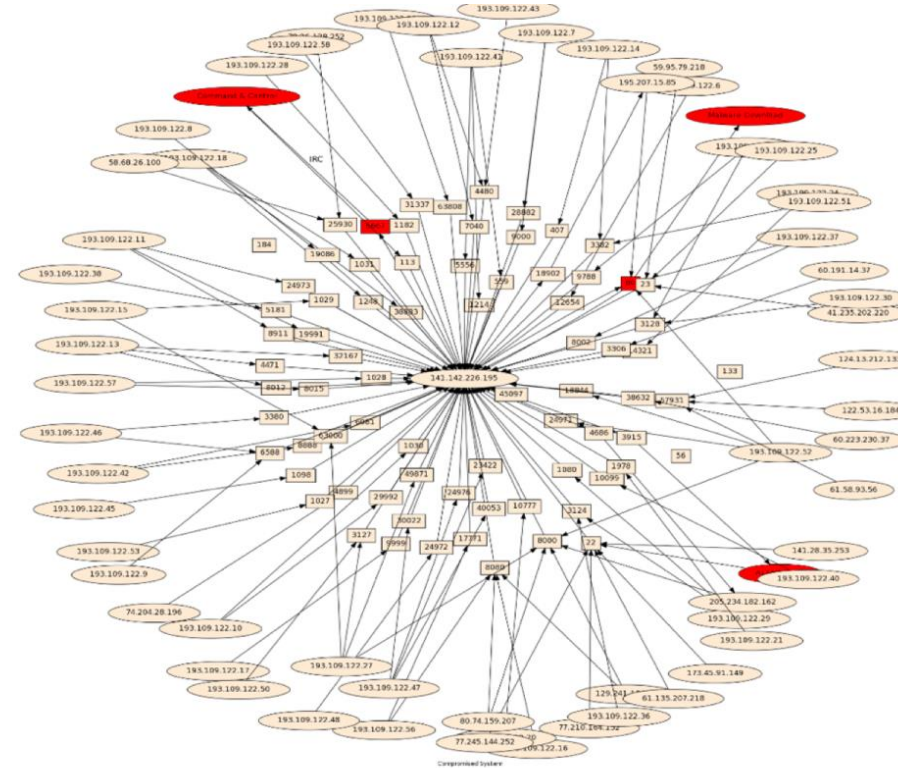


Five-Minute
Snapshot
of In-and-Out
Traffic
at NCSA

Five-Minute Snapshot of In-and-Out Traffic within NCSA@Illinois



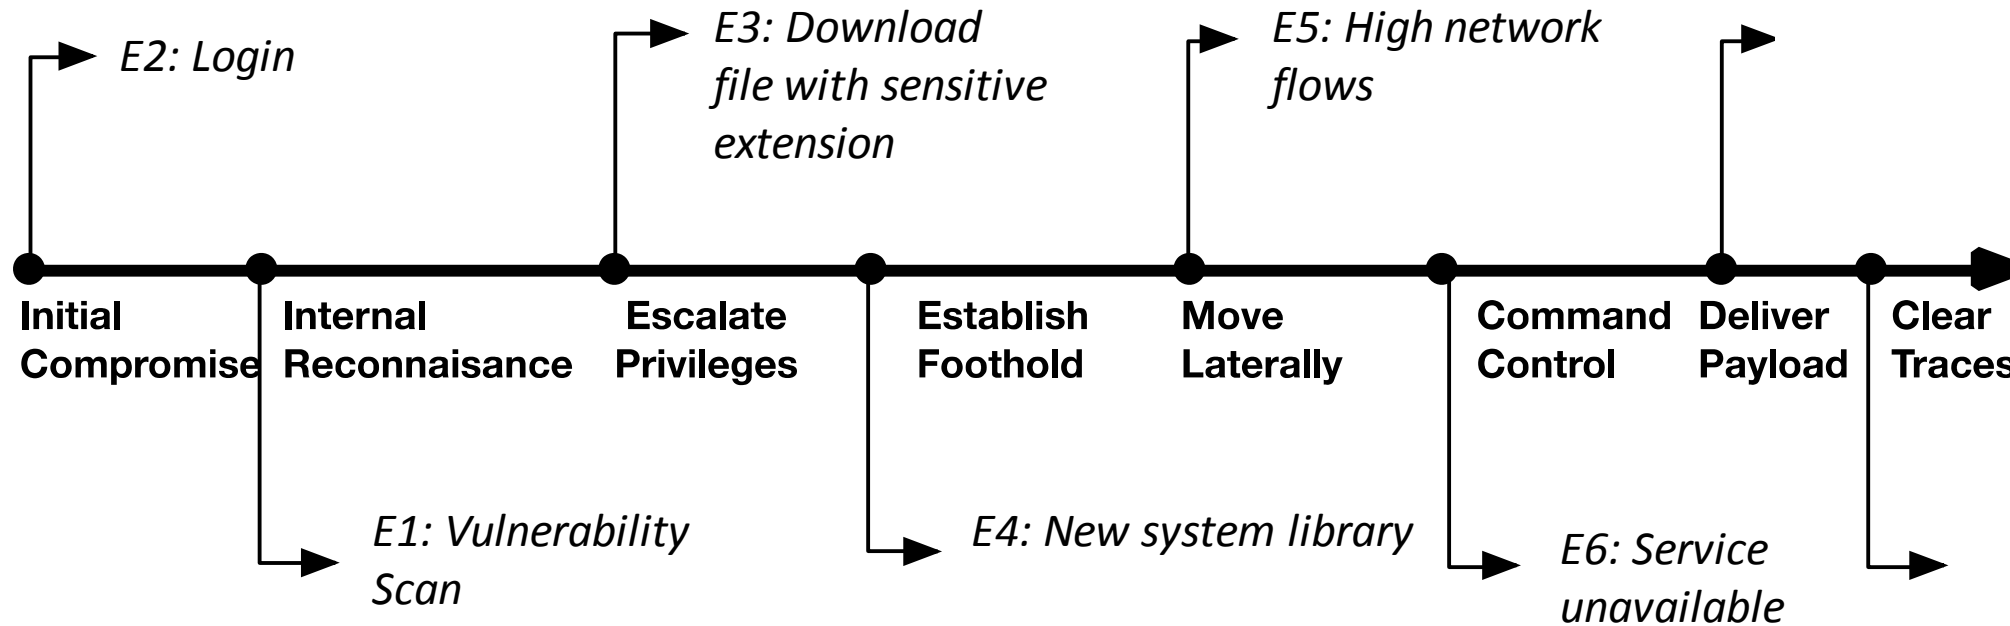
(a)



(b)

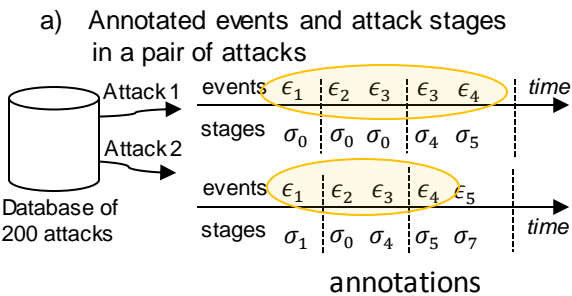
An Application in Security Data Analytics

Individual components of an attack as attack progresses



Attack stages for the credential stealing attack

Annotation and extracting patterns in past attacks



b) Event-stage annotation table for the attack pair (Attack 1 and Attack 2)

Event	Attack stage
$\{\epsilon_1\}$	$\{\sigma_0 \sigma_1\}$
$\{\epsilon_2\}$	$\{\sigma_0\}$
$\{\epsilon_3\}$	$\{\sigma_4\}$
$\{\epsilon_4\}$	$\{\sigma_5\}$
$\{\epsilon_5\}$	$\{\sigma_7\}$

c) Example patterns, stages, probabilities, and significance learned from the attack pair

Pattern	Attack stages	Probability in past attacks	Significance (p-value)
$[\epsilon_1, \epsilon_3, \epsilon_4]$	$[\sigma_1, \sigma_4, \sigma_5]$	q_a	p_a
$[\epsilon_1]$	$[\sigma_0 \sigma_1]$	q_b	p_b

...



Naïve Bayes

Bayesian Network

Dynamic Bayesian Network


Hidden Markov Model

Factor Graphs

OFFLINE ANNOTATION
ON PAST ATTACKS


OFFLINE LEARNING
OF PATTERNS

PROBABILISTIC GRAPHICAL MODELS



Note: ϵ_i is the corresponding value of an event E_t


Modeling the credential stealing attack using Naïve Bayes vs. Bayesian Network




Naïve Bayes

$$P(A|E_1, E_2, \dots, E_4) = P(A) \prod_i P(E_i|A)$$

Is (E1, E2, ..., E4) represents Benign activity?
 $[P(E_1|A = \text{Benign}) \dots P(E_4|A = \text{Benign})]P(A = \text{Benign}) > [P(E_1|A = \text{Attack}) \dots P(E_4|A = \text{Attack})]P(A = \text{Attack})$



$$P(E_1|A = 1)$$



Bayesian Network

Joint Distribution: $P(E_1, E_2, \dots, E_n, A) = P(A) \prod_{i=1}^n P(E_i|\text{parents}(E_i))$


Hypothesis:

$$P(A = \text{attack}|E_1, E_4, E_6) = ?$$


$$P(A = \text{benign}|E_1, E_4, E_6) = ?$$

ID	Description
A	Attack
E1	Vulnerability scan
E2	Login
E3	Download file with sensitive extension
E4	New system library
E5	High network flows
E6	Service unavailable

Modeling the credential stealing attack using Naïve Bayes vs. Bayesian Network



Naïve Bayes



Bayesian Network

ID	Description
A	Attack
E1	Vulnerability scan
E2	Login
E3	Download file with sensitive extension
E4	New system library
E5	High network flows
E6	Service unavailable

Model assumptions

1. All events share the same parent variable
2. All events are conditionally independent

Advantage:

Simplify calculation of posterior probability on A

Model assumptions

1. An event can be preceded (causal) by another event
2. There is no cycle in the network


Disadvantage

Explicitly assume causal relationships


(Causality may not be clear from the data)

For complicated attacks, causal loops may form and render the BN invalid


Modeling the credential stealing attack using Naïve Bayes vs. Bayesian Network




Naïve Bayes



Bayesian Network



Dynamic Bayesian Network



Hidden Markov Model

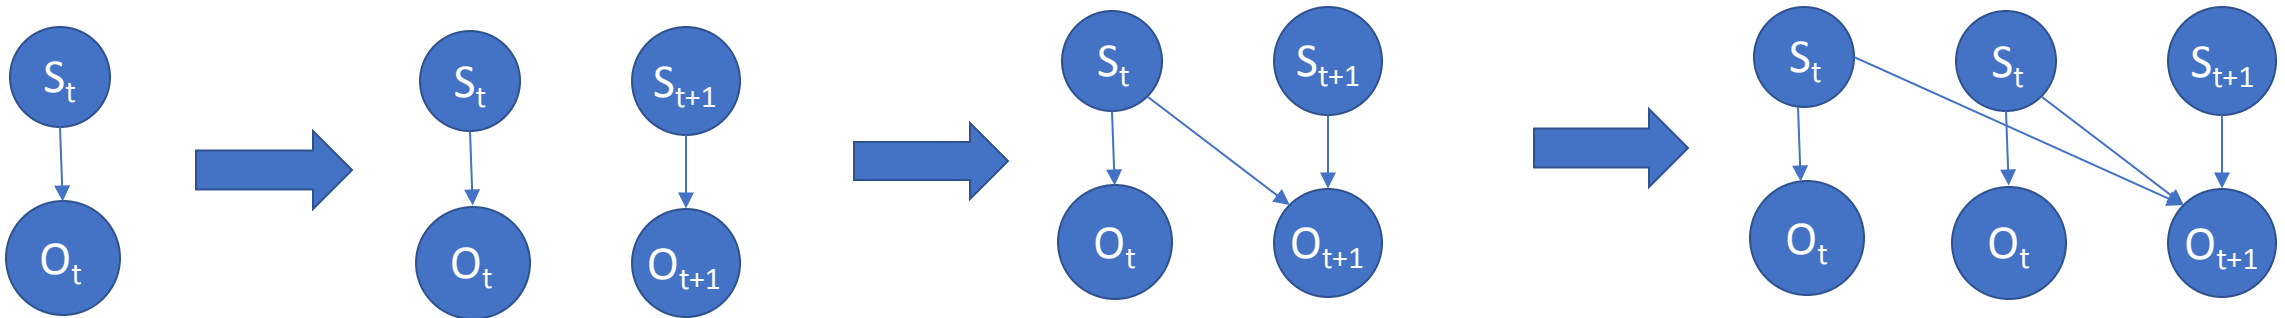
- When we consider the time evolution of the BN each variable in each timestep together, e.g., t and $t+1$, we have a Dynamic Bayesian Network that captures the first-order dependency --> referred to as the Markov Property
- This concept can be extended to higher order dependencies e.g on , $t-2$, $t-3$, ... and is called a higher-order Markov property, e.g., 2nd or 3rd Markov property.

$$P(A_1, E_1, \dots, A_n, E_n) = P(A_1)P(E_1|A_1) \dots P(E_{t+1}|A_{t+1})P(A_{t+1}|A_t)$$



Dynamic Bayesian Networks

- We have considered BNs with a static set of random variables, e.g., two variables: only one measurement variable and one state variable of the system.
- In reality, data is often time series in which each time step t has one measurement variable O_t and one state variable S_t . Thus, the number of random variables is proportional with the number of timesteps.
- Without correlating the random variables in each timestep, we have T disconnected BNs
- When we correlate each variable in each timestep together, e.g., t and $t+1$, we have a Dynamic Bayesian Network that captures the first-order Markov property.
- This concept can be extended for $t, t+1, t+2, \dots$ and is called a higher-order Markov property, e.g., 2nd or 3rd



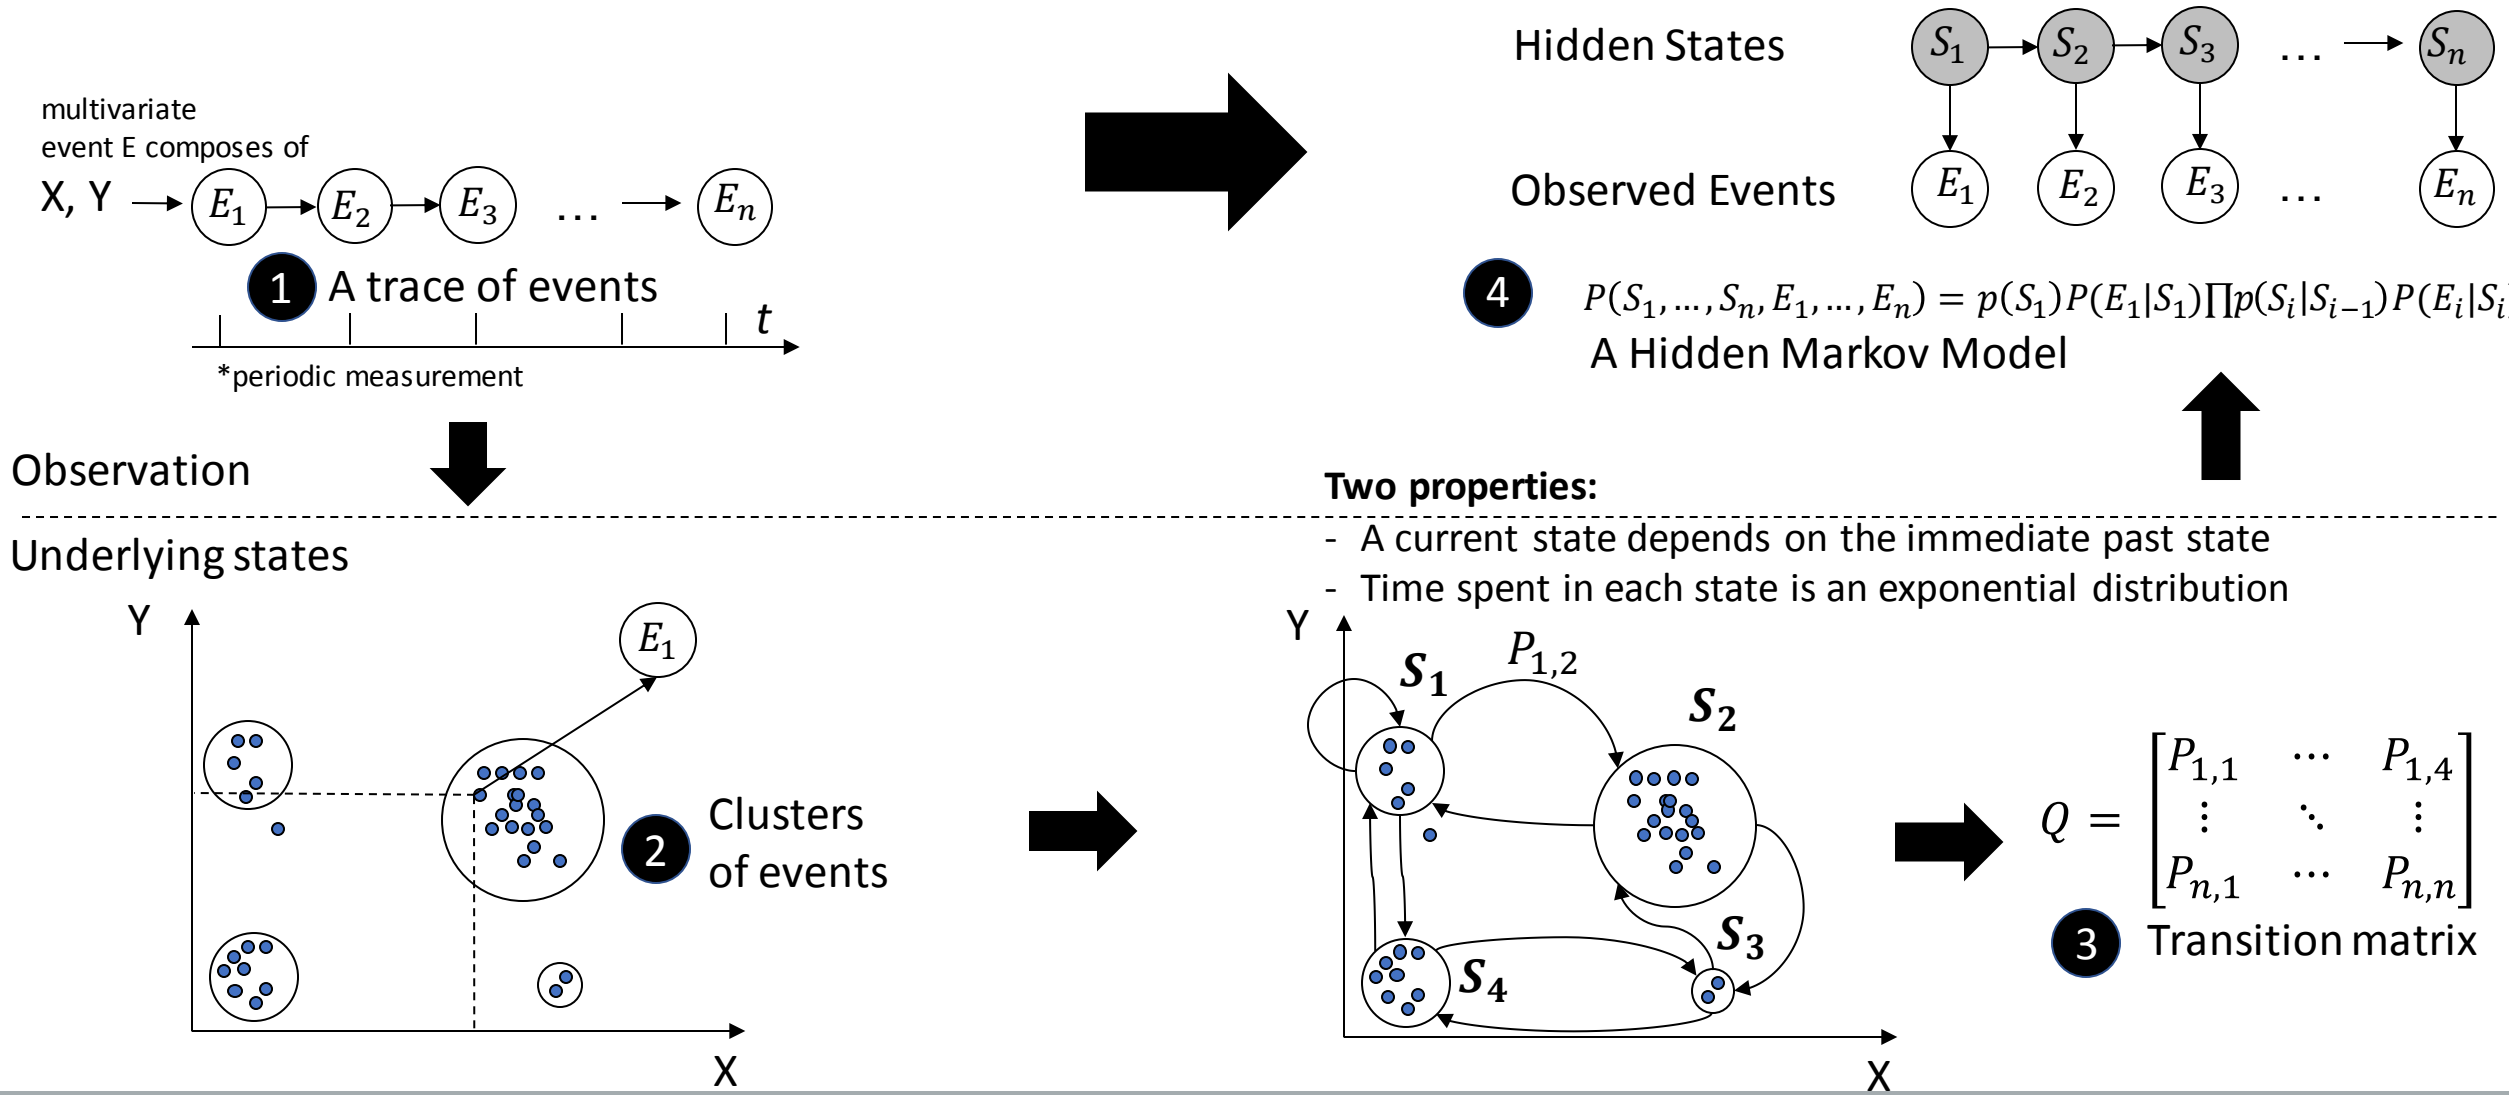
$$P(S_t, O_t) = P(S_t)P(O_t|S_t)$$

$$P(S_{t+1}, O_{t+1}) = P(S_{t+1})P(O_{t+1}|S_{t+1})$$

$$P(S_t, S_{t+1}, O_t, O_{t+1}) = P(S_t)P(O_t|S_t)P(O_{t+1}|S_t, S_{t+1})P(S_{t+1})$$

$$P(S_{t+1}, O_{t+1}) = P(S_{t+1})P(O_{t+1}|S_{t+1})$$

From a trace of events to a Hidden Markov Model



Hidden Markov Models

Model assumptions

An observation depends on its hidden state


A state variable only depends on the immediate previous state
(Markov assumption)

The future observations and the past observations are
conditionally independent given the current hidden state

Advantages:

HMM can model sequential nature of input data (future
depends on the past)

HMM has a linear-chain structure that clearly separates system
state and observed events.



$$P(S_1, \dots, S_n, E_1, \dots, E_n) = p(S_1)P(E_1|S_1)\prod p(S_i|S_{i-1})P(E_i|S_i)$$

**A Hidden Markov model on observed events
and system states**

Markov Model

- Consider a system which can occupy one of N discrete *states* or *categories*

$$x_t \in \{1, 2, \dots, N\} \longrightarrow \text{state at time } t$$



- We are interested in *stochastic* systems, in which state evolution is random
- Any *joint* distribution can be factored into a series of *conditional* distributions:

$$p(x_0, x_1, \dots, x_T) = p(x_0) \prod_{t=1}^T p(x_t \mid x_0, \dots, x_{t-1})$$

- For a *Markov* process, the next state depends only on the current state:

$$p(x_{t+1} \mid x_0, \dots, x_t) = p(x_{t+1} \mid x_t)$$


State Transition Diagrams



- Think of a particle randomly following an arrow at each discrete time step
- Most useful when N small, and Q *sparse*

Markov Chains: Graphical Models

$$p(x_0, x_1, \dots, x_T) = p(x_0) \prod_{t=1}^T p(x_t | x_{t-1})$$




Constraints on valid transition matrices:

$$q_{ij} \geq 0, \quad \sum_{i=1}^N q_{ij} = 1 \quad \text{for all } j$$

$$q_{ij} \triangleq p(x_{t+1} = i | x_t = j)$$

Hidden Markov Models

- Stall exists due to congestion
- Not directly measurable at runtime (hidden)
- Motivates *hidden Markov models* (HMM):



$$p(x_0, x_1, \dots, x_T) = p(x_0) \prod_{t=1}^T p(x_t | x_{t-1}) p(y_t | x_t)$$

Given x_t , previous observations impact future observations

$$p(y_t, y_{t+1}, \dots | x_t, y_{t-1}, y_{t-2}, \dots) = p(y_t, y_{t+1}, \dots | x_t)$$

State Transition Matrices

- A *stationary* Markov chain with N states is described by an $N \times N$ *transition matrix*:

$$Q = \begin{bmatrix} q_{11} & q_{12} & q_{13} \\ q_{21} & q_{22} & q_{23} \\ q_{31} & q_{32} & q_{33} \end{bmatrix}$$

$$q_{ij} \triangleq p(x_{t+1} = i \mid x_t = j)$$


- Constraints on valid transition matrices:

$$q_{ij} \geq 0 \quad \sum_{i=1}^N q_{ij} = 1 \quad \text{for all } j$$

State Transition Diagrams(Another Example)


$$q_{ij} \triangleq p(x_{t+1} = i \mid x_t = j)$$

$$Q = \begin{bmatrix} 0.5 & 0.1 & 0.0 \\ 0.3 & 0.0 & 0.4 \\ 0.2 & 0.9 & 0.6 \end{bmatrix}$$




- Think of a particle randomly following an arrow at each discrete time step
- Most interesting when Q *sparse*


Modeling the credential stealing attack using Naïve Bayes vs. Bayesian Network




Naïve Bayes



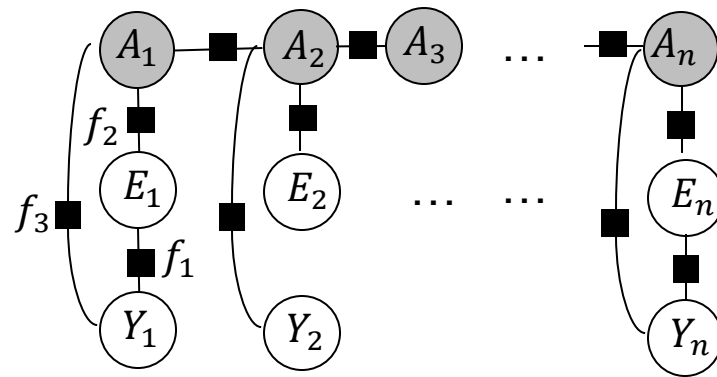
Bayesian Network




Dynamic Bayesian Network



Hidden Markov Model




Factor Graphs




Representing knowledge through graphical models

- A PGM encodes structural aspects of a joint probability distribution
 - $G = \{V, E\}$
- A node corresponds to a random variable
- An edge represent a dependencies between the variables



Why do we need graphical models?


- Graphs are an intuitive way of visualizing relationship among variables
- A graph shows the conditional independence between variables via edges
- Effective inference algorithms can be run on graphs such as belief propagation to infer marginal probabilities of variables




Definition of a Factor Graph

A factor graph is a **bipartite, undirected graph** of **random variables** and **factor functions**.
[Frey et. al. 01]


A factor function is a mathematical definition of ***prior beliefs*** or expert knowledge. *FG can represent both causal and non-causal relations.*




A factor graph for the product $f_A(x_1)f_B(x_2)f_C(x_1, x_2, x_3) \cdot f_D(x_3, x_4)f_E(x_3, x_5)$.




Bayesian Network (BN)



Factor Graph equivalent of BN



Markov Random Fields (MRF)



Factor Graph equivalent of MRF

Applications of Probabilistic Graphs in Security Domain

Problem statement. Given a set of security events, infer whether an attack is in progress?

Modeling Approach.

Each security event is a known variable \mathbf{e} , each takes value from a discrete set of events \mathbf{E} . An attack happens in a chain of exploits, thus we have a sequence of events in time dimension.

Each event is associated with a corresponding attack state \mathbf{s} , which is unknown. The simplest approach is to classify \mathbf{s} as a binary $\{0,1\}$. However, when we can infer \mathbf{s} it is often too late (the attacker is already in the system)

Thus, we want to discretize \mathbf{s} to smaller attack stages and provide update on such stages as soon as an event is observed.

Applications in the Security Domain (cont.)

Problem statement. Given a set of security events, infer whether an attack is in progress?

Formally, the problem becomes

1. Define a joint probability distribution function (joint pdf)

$$P(e_1, e_2, \dots, e_n, s_1, s_2, \dots, s_n)$$

2. Derive a conditional probability


$$P(e_1, e_2, \dots, e_n | s_1, s_2, \dots, s_n)$$

However, the search space is exponentially large (by the order of the number of observed stages and events) and the joint pdf is sophisticated.

We want to break the joint pdf into smaller components that are easier to compute, i.e., factorize the joint pdf.

Underlying representation of a Hidden Markov Model and conversion to a Factor Graph

Hidden Markov Model




Example

$|S| = 10$

$|X| = 10$


$|Y| = 10$



size of tensor

$10 \times 10 \times 10 = 1000$

Factor Graph of the HMM



Domain knowledge: "variables are pair-wise related" reduces dimensionality

size of three matrices + one transition


$10 \times 10 + 10 \times 10 + 10 \times 10 + 10 \times 10 = 400$

$1000^n \gg 400 \times n$

Modeling the credential stealing attack using Factor Graphs

OFFLINE ANNOTATION ON PAST ATTACKS

a) Annotated events and attack stages in a pair of attacks



b) Event-stage annotation table for the attack pair (Attack 1 and Attack 2)

Event	Attack stage
$\{\epsilon_1\}$	$\{\sigma_0 \sigma_1\}$
$\{\epsilon_2\}$	$\{\sigma_0\}$
$\{\epsilon_3\}$	$\{\sigma_4\}$
$\{\epsilon_4\}$	$\{\sigma_5\}$
$\{\epsilon_5\}$	$\{\sigma_7\}$

OFFLINE LEARNING OF PATTERNS

c) Example patterns, stages, probabilities, and significance learned from the attack pair

Pattern	Attack stages	Probability in past attacks	Significance (p-value)
$[\epsilon_1, \epsilon_3, \epsilon_4]$	$[\sigma_1, \sigma_4, \sigma_5]$	q_a	p_a
$[\epsilon_1]$	$[\sigma_0 \sigma_1]$	q_b	p_b

$$f(E) = \exp\{q_E(1 - p_E)\}$$

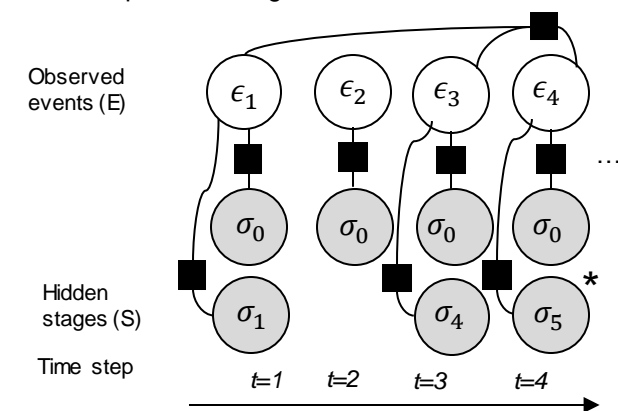
A factor function defined on the learned pattern, stages, and its significance

Model assumptions

1. There are multivariate relationships among the events
 2. Such relationships are represented by factor functions
 3. There is no restriction on order of the relationships like causal in Bayesian Network
- More suitable for modeling highly complex attacks, where the causal relations among the events are not immediately clear.


RUNTIME DETECTION OF UNSEEN ATTACKS

d) An evolution of the Factor Graph for the port knocking attack at run-time




Observed Security events	Factor function	ϵ_1 vulnerability scan	σ_0 benign
Unknown attack stages	Attack detected and stopped before the system misuse	ϵ_2 login	σ_1 discovery
		ϵ_3 sensitive_uri	σ_4 privilege escalation
		ϵ_4 new_library	σ_5 persistence


Taxonomy of graphical models




Naïve Bayes



Bayesian Network



Hidden Markov Model




Factor Graph

Conditional probabilities and statistical dependencies can be represented by a general type of graph: Factor Graph


Model structure and inference in PGMs

		Naïve Bayes	Bayesian Network	Hidden Markov Model	Factor Graphs
<i>MODEL STRUCTURE</i>	<i>Graph type</i>	Directed	Directed	Directed	Undirected
	<i>Graph structure</i>	Parent-child	Hierarchical parent-child	Sequential	Arbitrary structure
	<i>Variable of interest</i>	Attack (0 or 1)	Attack (0 or 1)	Sequence of system states	Sequence of attack stages
	<i>Relationship</i>	Conditional independence	Prior Conditional independence	State transitions Emission of event	Temporal relationships (patterns of events) Statistical relationships (severity or repetitiveness of events)
<i>INFERENCE</i>	<i>Algorithm</i>	Multiplication of conditional probabilities	Multiplication of conditional probabilities and priors	Dynamic Programming	Belief Propagation Sampling

Modeling the credential stealing attack using Naïve Bayes vs. Bayesian Network



Naïve Bayes



Bayesian Network

ID	Description
A	Attack
E1	Vulnerability scan
E2	Login
E3	Download file with sensitive extension
E4	New system library
E5	High network flows
E6	Service unavailable

Model assumptions

1. All events share the same parent variable
2. All events are conditionally independent

Advantage:

Simplify calculation of posterior probability on A

Model assumptions

1. An event can be preceded (causal) by another event
2. There is no cycle in the network

Disadvantage


Explicitly assume causal relationships

(Causality may not be clear from the data)

For complicated attacks, causal loops may form and render the BN invalid

Bayesian Networks vs. Markov Random Fields vs. Factor Graphs

Bayesian networks




$$p(x_1)p(x_2|x_1)p(x_3|x_1)$$

Product of
conditional
probabilities

Causal relationships

Markov random fields




$$\frac{1}{Z} \psi_1(x_1, x_2) \psi_2(x_1, x_3)$$

Product of
dependencies
among variable
cliques

Statistical dependencies

Factor graph

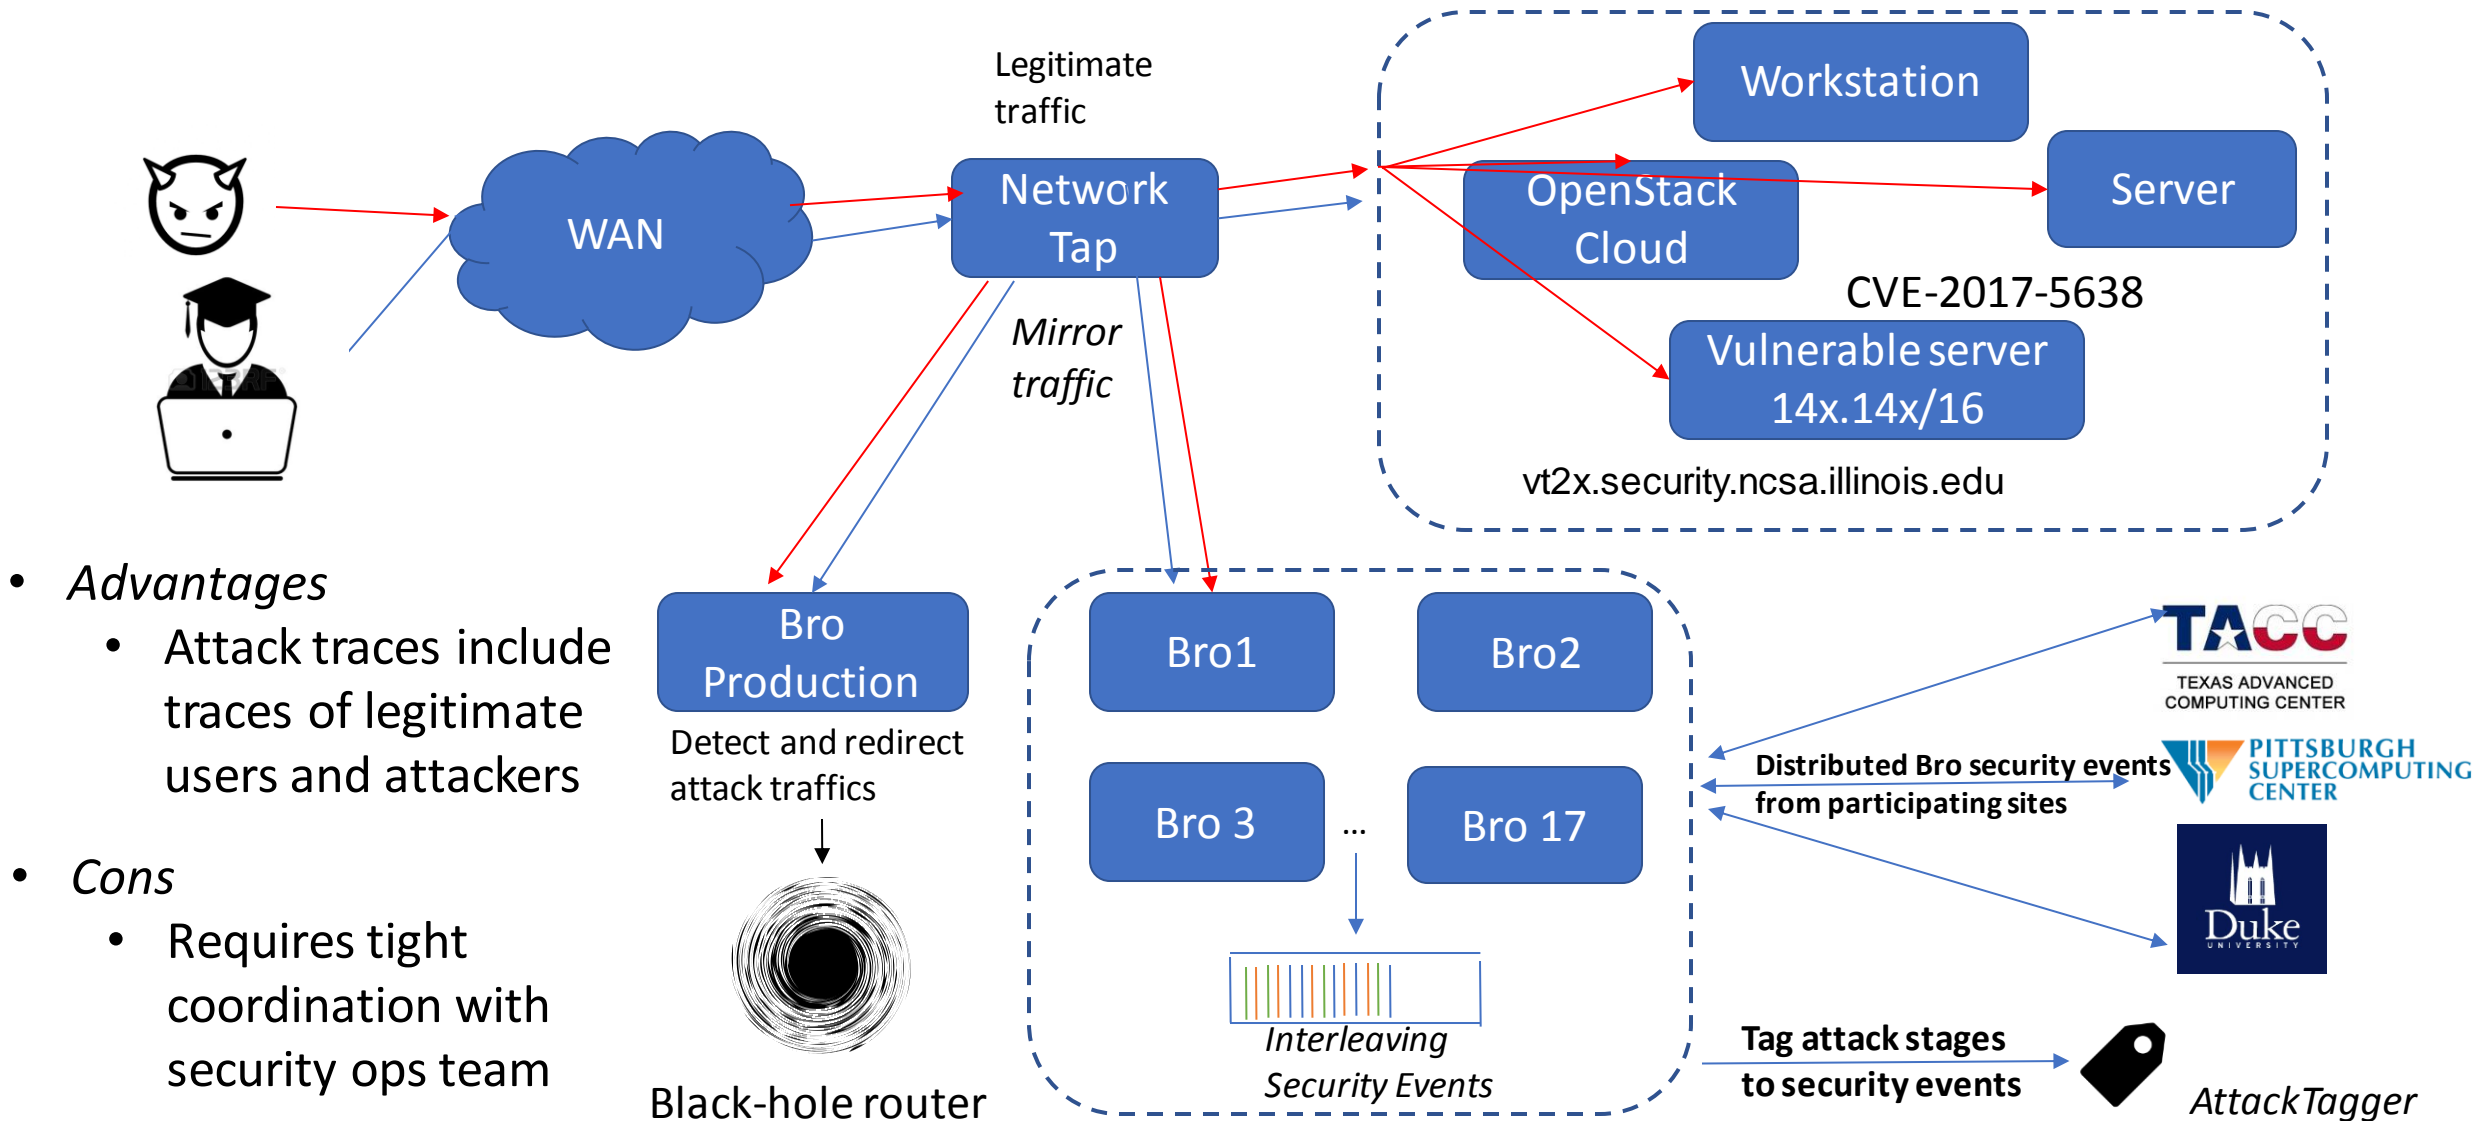


$$\frac{1}{Z} f_1(x_1) f_2(x_2, x_1) f_3(x_1, x_3)$$

Product of
dependencies using
univariate, bivariate, or
multivariate functions



Both types of relations
(including prior on a variable)

An attack testbed in real production traffic – an experiment at NCSA



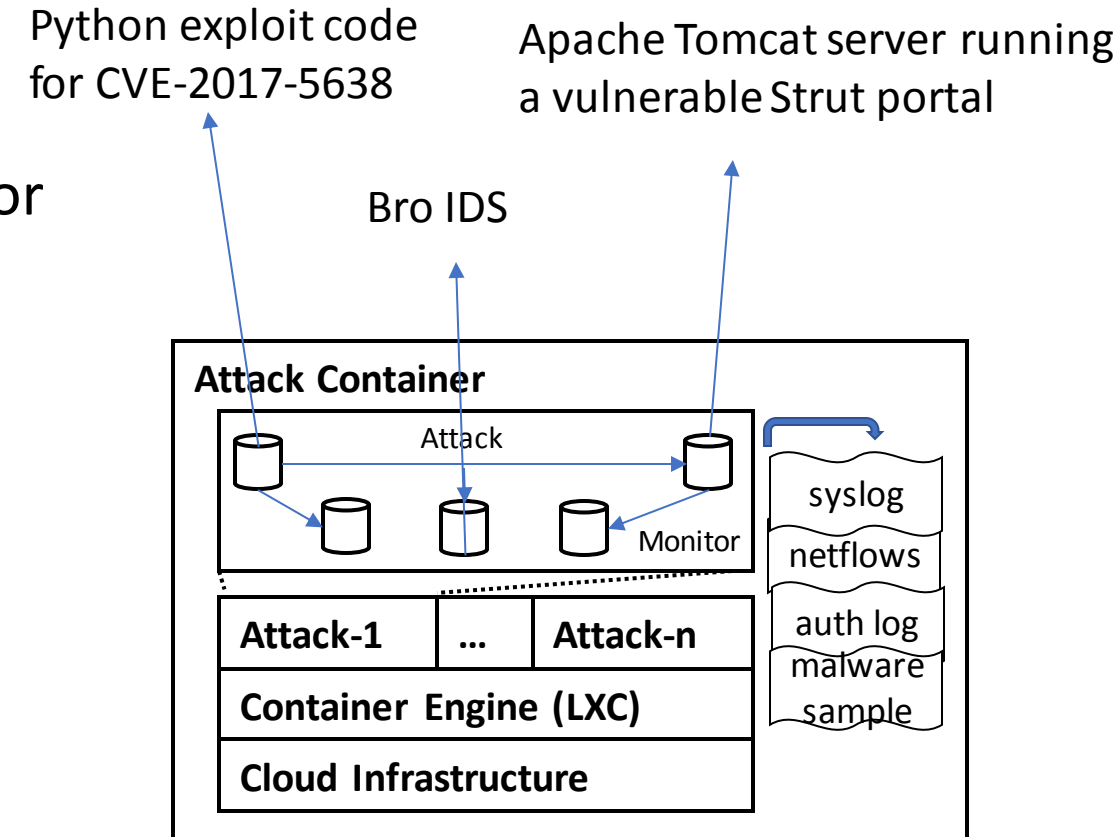
Evaluation Result

Stage transition of a multi-stage attack that exploits CVE-2017-5638

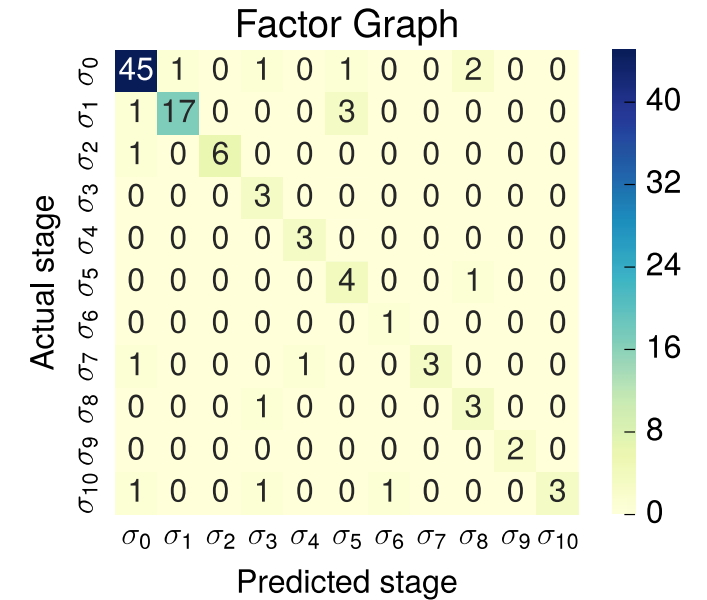




Emulating CVE-2017-5638 in a container-based environment

- *Advantages*
 - We were able to create an exact environment for the vulnerable Strut application
 - Monitors are in place to collect attack traces
 - Network policies are implemented to isolate potential outbreak of the attack
- *Limitations*
 - Containers are not exposed to a real network thus are not visible to attackers
 - Traces only include attack activities



Evaluation Results



Concluding Remarks

1. Probabilistic Graphical Models appear to be the way to integrate disparate issues on failure and attack pre-emption
2. Continuous and dynamic monitoring and adaptive abstraction offered by the factor graph based learning is critical
3. Going forward: Factor graphs could combine both security logs and error logs for diagnosis




6000+
users



5+ millions
connections



34M+
log events



4.5+ GB
Compressed final log

Heterogeneous host and network logs

Syslog
Netflows
IDS alerts
Human-written reports

200+ incidents in the past years
(2008-2017)


Brute-force attacks

Credential compromise

Abusing computing
infrastructure

Send spam

Launch Denial of
Service attacks.



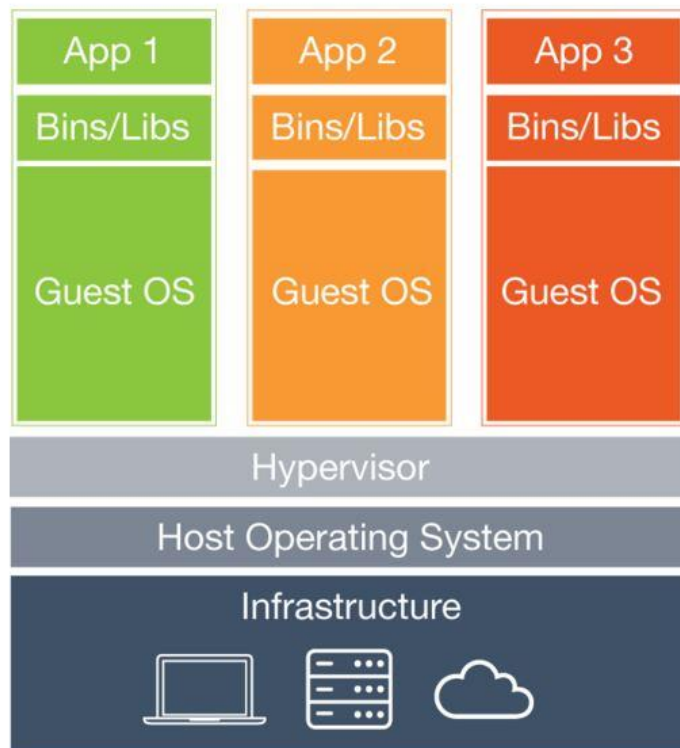
5-minute snapshot of network traffic in and out of
NCSA



Why attack injection?

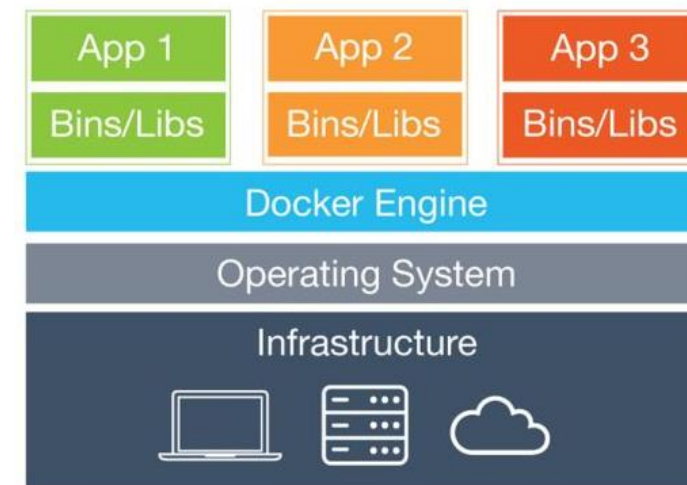
- Vulnerabilities are discovered on a daily basis, however, is a target system immune from such vulnerabilities?
- **Our goals are to:**
 - Evaluate ability of security monitoring systems in capturing attack-related security events
 - Run live, integration tests on applied security patches
 - Provide a dynamic blueprint of an attack (in terms of attack stages) as the attack unfolds across a production network

What is a Linux Container (LXC)?



Virtual Machine (VM) is an efficient, isolated duplicate of a real computer machine.

Features	Virtual Machine	Linux Container
Emulation	A real machine	A Linux system
Guest OS	Almost any OS	Only Linux system
Isolation and Resource management	Fully virtualized	Kernel namespace and control groups
File system	Separated file system for each VM	Layered filesystem (AUFS)
Disk and Memory	GBs	MBs
Startup time	Minutes	Seconds



Linux Container (LXC) is a virtualization technology for running multiple isolated Linux systems (containers).

How does AttackTagger work?


Dependencies Extraction

Commonality
Repetitiveness
Severity
....

User 1	-
User 2	+
...	...
User n	-


Ground truth on compromised users (+) and legitimate user (-)

Ranking and Selection



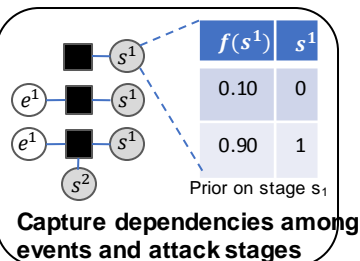
chkpwd[4495]: password failed
bro[820]: sensitive (venom.c)
ossec[918]: volatile /dev/shm
bro[820]: knocking SSH-2.5-OpenSSH_6.1.9

System logs



Raw logs of past attacks

Factor functions


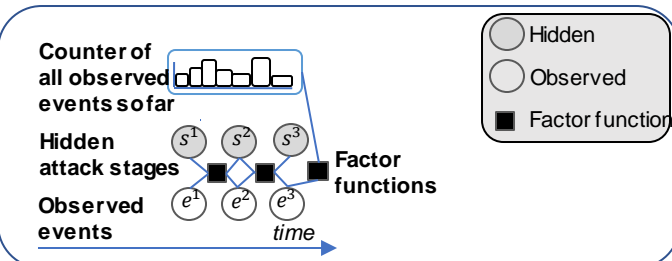


Capture dependencies among events and attack stages

Scan::Address_Scan 117.xxx.xxx.xxx scanned at least 7 unique hosts on port 22/tcp
SSL::Invalid_Server_Cert SSL certificate validation failed with (self signed certificate)


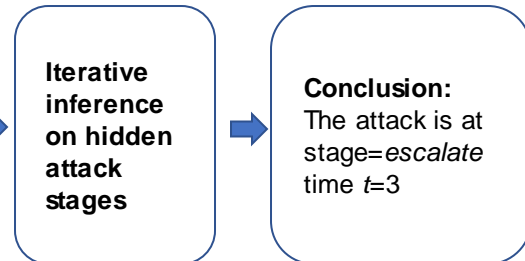
IDS alerts

Per-user factor graph



Use factor functions to connect selected events
... e^1 e^2 e^3
Timeline of observed security events

Predict an attack stage at every time step t



Automatic Learning of Factor Functions

Automatic Factor Graphs Generation

Runtime Attack Prediction

